

Advanced Inter-process Communications – Part1

Advanced Inter-process Communication

- Record locking with `fcntl`
- Three IPC facilities
 - Message passing
 - Allow a process to send and receive messages, a message is an arbitrary sequence of bytes or characters
 - Semaphores
 - Provide a rather low-level means of process synchronization, not suitable to the transmission of large amounts of information
 - Shared memory
 - Allow two or more processes to share the data contained in specific memory segments

Record Locking with fcntl

- `int fcntl(int fd, int cmd, struct flock *ldata);`
- Read lock
 - Simply prevents any process from applying the other type of fcntl lock, which is called a write lock.
 - Several processes may read lock the same segment simultaneously.
 - *fd* must have been opened using `O_RDONLY` or `O_RDWR`, so a file descriptor from `creat` won't work.
- Write lock
 - Stops any process from applying a read or write lock to the file.
 - *fd* must have been opened using `O_WRONLY` or `O_RDWR`.

File Lock Structure

- struct flock {
 short l_type; /* F_RDLCK, F_WRLCK, F_UNLCK */
 short l_whence; /* SEEK_SET, SEEK_CUR, SEEK_END */
 long l_start; /* byte offset, interpreted according to l_whence */
 long l_len; /* number of bytes to lock.
 long l_pid; /* ID of process that locked file */
};

File Lock Structure Fields

- `l_type`
 - `F_RDLCK` for read lock, `F_WRLCK` for write lock, and `F_UNLCK` for unlock operations.
- `l_whence`
 - `SEEK_SET`: (= 0) lock offset is from beginning of file.
 - `SEEK_CUR`: (= 1) current position of the file pointer.
 - `SEEK_END`: (= 2) end of file.
- `l_start`
 - Byte offset, interpreted according to `l_whence`.
- `l_len`
 - Number of bytes to lock
 - If 0, lock from `l_start` to end of file.

Lock Commands

- *cmd*
 - F_GETLK
 - Get lock description based on data passed via the *ldata* (describe the first lock that blocks the lock described in *ldata*)
 - F_SETLK
 - Lock or unlock the file as specified by *ldata*. Return *immediately* with -1 if unable to lock.
 - F_SETLKW
 - Lock or unlock data in file as specified by *ldata*. *Sleep* if unable to lock.
 - A process sleeping on a *fcntl* lock can be interrupted by a *signal*.

Features of Record Locking with fcntl()

- Locks are associated with a process and a file
- Locks are *never* inherited by the child across a fork()
- A call to fcntl() does not alter the file's read-writer pointer
- All locks belonging to a process are removed automatically when the process dies

Example #1: Record Lock (1)

```
#include <fcntl.h>
#include <unistd.h>
#include <stdlib.h>

main()
{
    int fd;
    struct flock my_lock;

    /* 쓰기 록의 인수를 지정 */
    my_lock.l_type = F_WRLCK;
    my_lock.l_whence = SEEK_SET;
    my_lock.l_start = 0;
    my_lock.l_len = 10;

    /* 파일을 개방한다. */
    fd = open ("locktest", O_RDWR);
```


Example #1: Record Lock (2)

```
/* 처음 10바이트를 록한다 */
if (fcntl (fd, F_SETLKW, &my_lock) == -1)
{
    perror ("parent: locking");
    exit (1);
}
printf ("parent: locked record\n");
switch (fork()){
case -1:                /* 오류 */
    perror ("fork");
    exit (1);
case 0:                /* 자식 */
    my_lock.l_len = 5;
    if (fcntl (fd, F_SETLKW, &my_lock) == -1)
    {
        perror ("child: locking");
        exit (1);
    }
    printf ("child: locked\n");
    printf ("child: exiting\n");
    exit(0);
}
```

Example #1: Record Lock (3)

```
sleep(5);
```

```
/* 이제 퇴장(exit)한다. 따라서 록이 해제된다 */  
printf ("parent: exiting\n");  
exit (0);  
}
```

```
dhlee@kde:~/Course/SP/example>flick-test
```

```
parent: locked record
```


```
parent: exiting
```

```
child: locked
```

```
child: exiting
```

Dead Lock 복습

- Dead Lock의 국문용어?
- Dead Lock이 일어날 4가지 조건
 - Mutual exclusion
 - Hold and wait
 - No preemption
 - Circular wait
- Dead Lock Handling
 - Prevention
 - Avoidance
 - Detection & Recovery



RAG

Example #2 : Dead Lock (1)

```
#include <fcntl.h>
#include <unistd.h>
#include <stdlib.h>

main()
{
    int fd;
    struct flock first_lock;
    struct flock second_lock;

    first_lock.l_type = F_WRLCK;
    first_lock.l_whence = SEEK_SET;
    first_lock.l_start = 0;
    first_lock.l_len = 10;

    second_lock.l_type = F_WRLCK;
    second_lock.l_whence = SEEK_SET;
    second_lock.l_start = 10;
    second_lock.l_len = 5;
```

Example #2 : Dead Lock (2)

```
fd = open ("locktest", 0_RDWR);

if ( fcntl (fd, F_SETLKW, &first_lock) == -1)      /* A */
    fatal ("A");

printf ("A: lock succeeded (proc %d)□n", getpid());

switch (fork()){
case -1:
    /* 오류 */
    fatal ("error on fork");
case 0:
    /* 자식 */
    if ( fcntl (fd, F_SETLKW, &second_lock) == -1) /* B */
        fatal ("B");
    printf ("B: lock succeeded (proc %d)□n", getpid());

    if ( fcntl (fd, F_SETLKW, &first_lock) == -1) /* C */
        fatal ("C");
    printf ("C: lock succeeded (proc %d)□n", getpid());
    exit (0);
```

Example #2 : Dead Lock (3)

default:

```
    /* 부모 */  
    printf ("parent sleeping\n");  
    sleep (10);  
    if ( fcntl (fd, F_SETLKW, &second_lock) == -1) /* D */  
        fatal ("D");  
    printf ("D: lock succeeded (proc %d)\n", getpid());  
}  
}
```

dhlee@kde:~/Course/SP/example>deadlock

A: lock succeeded (proc 21884)

B: lock succeeded (proc 21885)

parent sleeping

D: Resource deadlock avoided

D: lock succeeded (proc 21884)

C: lock succeeded (proc 21885)

IPC Common Features (1)

- Three commonly used IPC facilities
 - Message passing
 - Semaphores
 - Shared Memory
- An IPC facility must be created prior to use.
 - It can be created anytime before use.
- Permission and ownership for a facility are established at the time of creation.
 - But these attributes can be changed later.
- An IPC facility can exist with contents intact after all processes accessing it have exited.
 - Changes made to an IPC facility by a process persist after the process exits.
- An IPC facility must be explicitly removed.
 - Removal of an IPC facility can be done by either the owner or creator.

IPC Common Features (2)

- Needs an IPC *key* to gain access to an IPC facility.
- All processes wishing to use the same IPC facility must specify the same key.
 - This key must uniquely identify the IPC facility.
- An IPC facility is identified by a nonnegative integer.
 - In need of an IPC key to generate such a numeric ID.
- Three ways to generate an IPC key:
 - Let the system pick a key (IPC_PRIVATE)
 - Pick a key directly (use a predefined key and store it in a common header file)
 - Use `key_t ftok(const char *path, int id)` to generate an IPC key.

Message Queues

- A message queue allows processes to send and receive discrete amounts of data, known as *messages*.
 - A message is a collection of bytes of varying lengths, ranging from zero to a system imposed maximum.
 - The message content can be anything -- ASCII or binary.
- Each message has a type associated with it.
 - A message type is a long integer.

Semaphores

- Semaphores are useful for process synchronization and resource management.
 - A semaphore is an unsigned integer shared among several competing processes.
 - Operating upon a semaphore involves adding or subtracting from the underlying semaphore value.
- The UNIX implementation of semaphores is modeled after Dijkstra's semaphores.

Shared Memory

- Shared memory allows one or more process to share the same segment of memory.
 - Shared memory may be attached to more than one process.
 - Once attached, it is part of a process's data space.
- Shared memory is the fastest IPC.

IPC System Call Categories

- *get*: Is used to create or access an IPC facility. It returns an IPC identifier.
 - `semget()`, `shmget()`, `msgget()`
- *ctl*: Is used to determine status, set options and/or permissions, or remove an IPC facility.
 - `semctl()`, `shmctl()`, `msgctl()`
- *op*: Is used to operate on an IPC identifier.
 - `semop()`
 - `shmat()`, `shmdt()`
 - `msgsnd()`, `msgrcv()`

End-User Level Commands

- `ipcs`
 - Report inter-process communication facilities status.
- `ipcrm`
 - Remove a message queue, semaphore set, or shared memory ID.

Message Queues (1)

- Message queues are for exchange of data among processes.
- The data exchanged are in discrete portions known as *messages*.
 - A message is a sequence of bytes, from zero to the system limit, with a message type.
 - The structure of a message is up to the programmer.
- Once a message queue is created, processes can send/receive messages to/from the queue.
 - A process can selectively read messages on a queue:
 - First message on the queue.
 - First message of a specific type on the queue.
 - First message from a range of types (from 1 to upper bound of range) on the queue.

Message Queues (2)

- Messages queues are similar to unnamed pipes; however,
 - Each message has a type associated with it.
 - Data in a pipe is just a sequence of types with no header of format.
 - Unread data in a pipe vanishes when the last reader process close a file descriptor associated with the pipe's read end.
 - A message queue sticks around.
 - It will remain active until it is explicitly removed.

Message System Calls

- `msgget()`
 - To create or to gain access to a message queue.
- `msgctl()`
 - To determine the status of a message queue.
 - To change the permission or ownership of a message queue.
 - To change the maximum size.
 - To remove a message queue.
- `msgsnd()`
 - To send a message.
- `msgrcv()`
 - To receive a message.

Some Comments

- The System V IPC structures are system-wide.
 - They stick around until explicitly removed.
- Not compatible with standard UNIX file manipulation facilities.
 - Can't use `open()`, `close()`, `read()`, `write()`, etc.
 - Almost a dozen of new system calls were added to the kernel to support these IPC facilities.
- Basically, they are intra-system IPC facilities.

msgget

- `#include <sys/types.h>`
`#include <sys/ipc.h>`
`#include <sys/msg.h>`
`int msgget(key_t key, int flag);`
 - `msgget` returns an identifier to a message queue whose name is *key*.
 - *key* can specify that the returned queue identifier should refer to a private queue (`IPC_PRIVATE`), in which case a new message queue is created.
 - *flag* specifies if the queue should be created (`IPC_CREAT`), and if creation of the queue should be exclusive (`IPC_EXCL`). In the latter case, `msgget` fails (-1 will be returned) if the queue already exists.
 - `mgid = msgget((key_t)0100, 0644|IPC_CREAT|IPC_EXECL)`

msgsnd and msgrcv (1)

- `#include <sys/types.h>`
`#include <sys/ipc.h>`
`#include <sys/msg.h>`
`int msgsnd(int id, struct msgbuf *msgp, int size, int flag);`
`int msgrcv(int id, msgbuf *msgp, int size, int type, int flag);`
 - msgsnd sends a message of *size* bytes in the buffer *msgp* to the message queue *id*. msgbuf is defined as
 - `struct msgbuf {`
`long mtype;`
`char mtext[];`
`};`

msgsnd and msgrcv (2)

- If the IPC_NOWAIT bit is off in *flag*, msgsnd sleeps if the number of bytes on the message queue exceeds the maximum, or if the number of system-wide messages exceeds a maximum value. If IPC_NOWAIT is set, msgsnd returns immediately in these cases
- msgrcv receives messages from the queue identified by *id*.
 - If *type* is 0, the first message on the queue is received;
 - If positive, the first message of that type is received;
 - If negative, the first message of the lowest type less than or equal to *type*'s absolute value is received.
 - Example
 - The queue contains three message with mtype value of 999, 5, 1. if type of msgrcv is 0, returns 999; if 5, returns 5; if -999, returns in order 1, 5, and 999;

msgsnd and msgrcv (3)

- *size* in `msgrcv` indicates the maximum size of message text the user wants to receive. If `MSG_NOERROR` is set in *flag*, the kernel truncates the received message if its size is larger than *size*. Otherwise it returns an error.
- If `IPC_NOWAIT` is not set in *flag*, `msgrcv` sleeps until a message that satisfies *type* is sent. If `IPC_NOWAIT` is set, it returns immediately.
- `msgrcv` returns the number of bytes in the message text.

Example #3: Message (1)

```
/* q.h -- header for message facility example */

#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/msg.h>
#include <errno.h>

extern int errno;

#define QKEY (key_t)0150 /*identifying key for queue*/
#define QPERM 0660 /*permissions for queue*/
#define MAXOBN 50 /*maximum length obj. name*/
#define MAXPRIOR 10 /*maximum priority level*/

struct q_entry { /*struct we well use for message*/
    long mtype;
    char mtext[MAXOBN+1];
};
```

Example #3: Message (2)

```
/* function prototypes */  
int enter(char *objname, int priority);  
int init_queue(void);  
int serve(void);
```

Example #3: Message (3)

```
/* enter -- place an object into queue */

#include <string.h>
#include "q.h"

static int s_qid = -1; /*message queue identifier*/

int enter(char *objname, int priority)
{
    int len;
    struct q_entry s_entry; /*structure to hold message*/

    /*validate name length, priority level*/
    if ((len = strlen(objname)) > MAXOBN) {
        fprintf(stderr, "name too logn");
        return (-1);
    }
}
```


Example #3: Message (4)

```
if (priority > MAXPRIOR || priority < 0) {
    fprintf(stderr, "invalid priority level");
    return (-1);
}

/*initial message queue as necessary*/
if (s_qid == -1 && (s_qid = init_queue()) == -1)
    return (-1);

/* initialize s_entry */
s_entry.mtype = (long) priority;
strncpy(s_entry.mtext, objname, MAXOBN);

/*send message, waiting if necessary*/
if (msgsnd(s_qid, &s_entry, len, 0) == -1) {
    perror("msgsnd failed");
    return (-1);
}
else
    return (0);
}
```

Example #3: Message (5)

```
/* init_queue -- get queue identifier */

#include "q.h"

int init_queue(void)
{
    int queue_id;

    /*attemp to create message queue*/
    if ((queue_id = msgget(QKEY, IPC_CREAT|QPERM)) == -1)
        perror("msgget failed");
    return (queue_id);
}
```

Example #3: Message (6)

```
/*server -- serve object with highest priority on queue */

#include "q.h"

static int r_qid = -1;

Int proc_obj(struct q_entry *msg)
{
    printf("\npriority: %ld name: %s\n", msg->mtype, msg->mtext);
}

int serve(void)
{
    struct q_entry r_entry;
    int mlen;

    /*initialize queue as necessary*/
    if (r_qid == -1 && (r_qid = init_queue()) == -1)
        return (-1);
}
```

Example #3: Message (7)

```
/*get and process next message, waiting if necessary*/
for(;;) {
    if ((mlen = msgrcv(r_qid, &r_entry, MAXOBN, (long) -1*MAXPRIOR,
                      MSG_NOERROR)) == -1) {
        perror("msgrcv failed");
        return (-1);
    }
    else {
        /*make sure we've a string*/
        r_entry.mtext[mlen] = '\0';

        /*print object name*/
        proc_obj(&r_entry);
    }
}
}
```

Example #3: Message (8)

```
/*etest -- enter object name*/
```

```
#include <stdio.h>
```

```
#include "q.h"
```

```
int main(int argc, char *argv[])
```

```
{
```

```
    int priority;
```

```
    if (argc != 3) {
```

```
        fprintf(stderr, "usage: %s objname priority\n", argv[0]);
```

```
        exit(1);
```

```
    }
```

```
    if ((priority = atoi(argv[2])) <= 0 || priority > MAXPRIOR) {
```

```
        fprintf(stderr, "invalid priority");
```

```
        exit(2);
```

```
    }
```

Example #3: Message (9)

```
if (enter(argv[1], priority) < 0) {  
    fprintf(stderr, "enter failure");  
    exit(3);  
}  
exit(0);  
}
```

Example #3: Message (10)

```
/*stest -- simple server for queue*/

#include <stdio.h>
#include "q.h"

int main(void)
{
    int pid;

    switch(pid = fork()) {
        case 0: /*child*/
            serve();
            break; /*actually, serve never exits */
        case -1:
            perror("fork to start sever failed");
            break;
        default:
            printf("server process pid is %d\n", pid);
    }
    exit(pid == -1 ? 1 : 0);
}
```

Example #3: Sample Run

```
dhlee@kde:~/Course/SP/example>etest objname1 3
dhlee@kde:~/Course/SP/example>etest objname2 4
dhlee@kde:~/Course/SP/example>etest objname3 1
dhlee@kde:~/Course/SP/example>etest objname4 9
dhlee@kde:~/Course/SP/example>stest
```

priority: 1 name: objname3

priority: 3 name: objname1

priority: 4 name: objname2

priority: 9 name: objname4

server process pid is 22496

msgctl (1)

- `#include <sys/types.h>`
`#include <sys/ipc.h>`
`#include <sys/msg.h>`
`int msgctl(int id, int cmd, struct msqid_ds *buf);`
 - `msgctl` allows process to set or query the status of the message queue *id*, or to remove the queue, according to the value of *cmd*.

msgctl (2)



– The structure `msqid_ds` is defined as follows:

- `struct msqid_ds {`
 `struct ipc_perm msg_perm; /*permission struct */`
 `ushort msg_qnum; /*number of message on q*/`
 `ushort msg_qbytes; /*max number of bytes on q*/`
 `ushort msg_lspid; /*pid of last msgsnd operation*/`
 `ushort msg_lrpid; /*pid of last msgrcv operation*/`
 `time_t msg_stime; /*last msgsnd time*/`
 `time_t msg_rtime; /*last msgrcv time*/`
 `time_t msg_ctime; /*last change time*/`
};

msgctl (3)

- The commands (*cmd*) and their meaning are as follows:
 - IPC_STAT
 - Read the message queue header associated with *id* into *buf*.
 - IPC_SET
 - Set the values of `msg_perm.uid`, `msg_perm.gid`, `msg_perm.mode` (9 low-order bits), and `msg_qbytes` from the corresponding values in *buf*.
 - IPC_RMID
 - Remove the message queue for *id*.

Example #4: Message (1)

```
/* showmsg -- show message queue details */
```

```
#include <sys/types.h>
```

```
#include <sys/ipc.h>
```

```
#include <sys/msg.h>
```

```
#include <time.h>
```

```
#include <stdio.h>
```

```
void mqstat_print(int mqid, struct msqid_ds *mstat)
```

```
{
```

```
    printf("\nmsg_qid %d\n\n", mkey, mqid);
```

```
    printf("%d message(s) on queue\n\n", mstat->msg_qnum);
```

```
    printf("Last send by proc %d at %s\n", mstat->msg_lspid, ctime(&(mstat->msg_stime)));
```

```
    printf("Last recv by proc %d at %s\n", mstat->msg_lrpid, ctime(&(mstat->msg_stime)));
```

```
}
```

Example #4: Message (2)

```
int main(int argc, char *argv[])
{
    int msq_id;
    struct msqid_ds msq_status;

    if(argc != 2) {
        fprintf(stderr, "usage:showmsg keyval\n");
        exit(1);
    }

    /*get message queue identifier*/
    msq_id = atoi(argv[1]);

    /*get status information*/
    if (msgctl(msq_id, IPC_STAT, &msq_status) < 0) {
        perror("msgctl failed");
        exit(3);
    }
}
```

Example #4: Message (3)

```
/*print out status information*/  
mqstat_print(msq_id, &msq_status);  
exit(0);  
}
```

Example #4: Sample Run

```
dhlee@kde:~/Course/SP/example>ipcs
```

```
----- Shared Memory Segments -----
```

key	shmid	owner	perms	bytes	nattch	status
0x00000000	65537	root	777	393216	2	dest
0x00000000	98306	root	644	106496	3	dest
0x00000000	131075	root	644	106496	3	dest

```
----- Semaphore Arrays -----
```

key	semid	owner	perms	nsems
-----	-------	-------	-------	-------

```
----- Message Queues -----
```

key	msqid	owner	perms	used-bytes	messages
0x00000068	163840	dhlee	660	0	0

```
dhlee@kde:~/Course/SP/example>showmsg 163840
```

```
msg_qid 163840
```

```
0 message(s) on queue
```

```
Last send by proc 22494 at Tue Nov 16 11:27:35 2004
```

```
Last recv by proc 22496 at Tue Nov 16 11:27:35 2004
```

```
dhlee@kde:~/Course/SP/example>ipcrm msg 163840
```

```
msgrcv failed: Identifier removed  
resource(s) deleted
```